

Terra Training Session May 30th, 2023

**Welcome to a Terra Training Session hosted by Theiagen Genomics & the
Massachusetts Department of Public Health**

We appreciate your punctuality! Please give others a few minutes to arrive (and adjust their audio equipment). We will get started at **2:33 PM Eastern Time**. Thanks!





Docker for Public Health Bioinformatics

Week 2 - Container Repositories and Writing Dockerfiles

Tuesday May 30th, 2023

Curtis Kapsak, MS & Frank Ambrosio, MS | Theiagen Genomics

Course Introduction

Training Workshop Overview

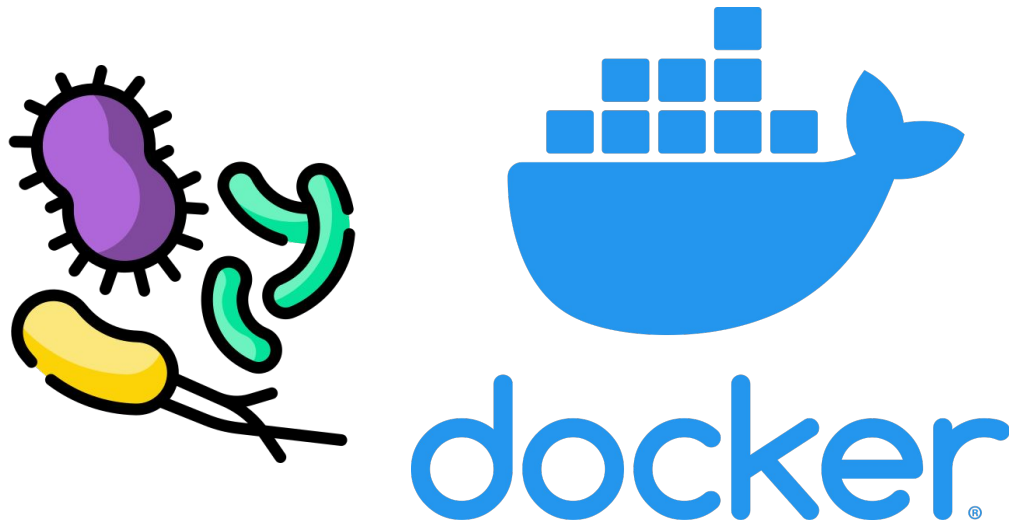
Training Information, Communication, and Support

- [Training Notion Page](#) created to host training resources and information
- **Support Contacts:**
 - support@terrapublichealth.zendesk.com



Main Course Objective

Learn about the concepts of Docker & containerization and their applications in public health bioinformatics



Training Workshop Overview

This workshop is an Intermediate/Advanced course

Great resources for more information regarding containers and pathogen genomics

For more **technical content**, get connected with various **pathogen genomics communities** such as PHA4GE, StaPH-B, & micro-binfie

- [StaPH-B Docker User Guide](#)
- [Ten Recommendations for supporting open pathogen genomic analysis in public health](#)
 - Highlights containers and workflow management systems in context of public health
- [A Primer on Infectious Disease Bacterial Genomics](#)
 - Introduction to analyzing pathogen genomics data

Course Structure



4-Week Virtual Training Workshop

- **All training sessions** will begin at 2:30pm Eastern Time
 - Live Lectures (90m) on Mondays
 - Office Hours (60m) on Wednesdays
 - Exceptions:
 - No sessions the week of APHL Annual conf. (May 22-25)
 - Week 2 lecture will occur Tue May 30th 2:30-4pm EST due to Memorial Day
- **Live lectures** will include **hands-on exercises**
 - To participate, please ensure that you have registered for a GitHub account

Course Content

Week One - Intro to Docker and Containerization

- **Lecture Content:** Introduction to Docker containers
- **Hands-on Exercises:** Utilize a docker container to download a *Klebsiella pneumoniae* genome and to run Kleborate

Week Two - Container Repositories and Writing Dockerfiles

- **Lecture Content:** Intro to various repositories for Docker containers e.g. StaPH-B docker-builds & biocontainers
- **Hands-On Exercise:** Build docker images using pre-existing dockerfiles

Course Content

Week Three - Developing custom Docker Images

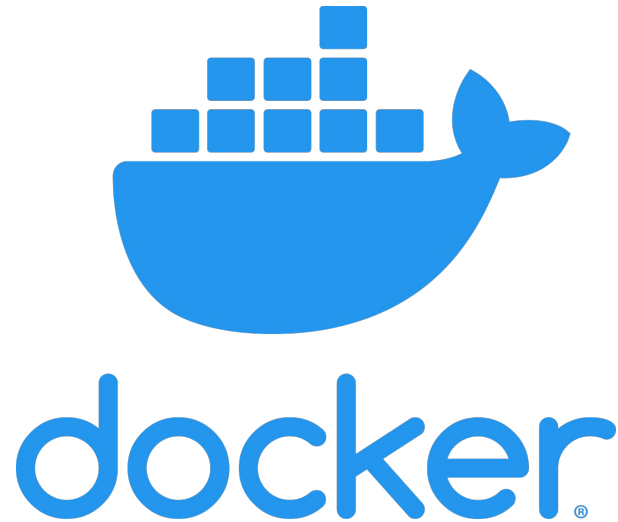
- **Lecture Content:** Intro to development and testing practices for writing dockerfiles
- **Hands-on Exercise:** TBD

Week Four - StaPH-B docker-builds project

- **Lecture Content:** Review of the StaPH-B docker-builds project and code repository
- **Hands-On Exercise:** Develop a dockerfile and create a pull request

Week 2

Container Repositories and Writing Dockerfiles

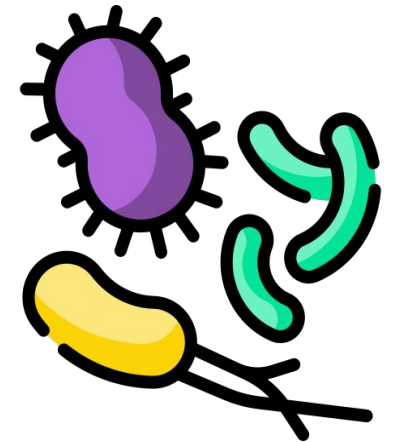


Container Repositories and Writing Dockerfiles

Lecture Content: Intro to various repositories for public health related Docker images and writing Dockerfiles

Hands-on Exercises:

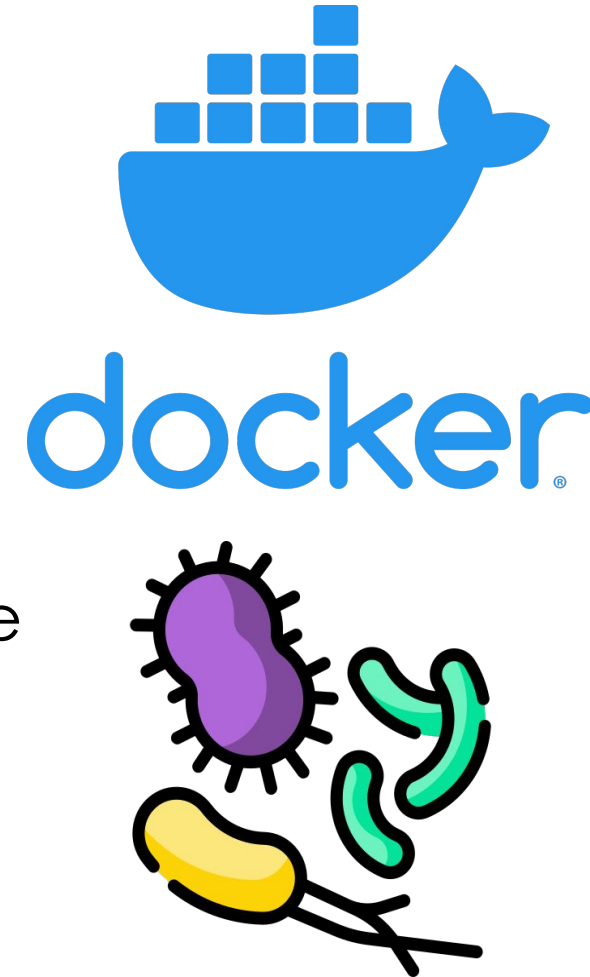
- **Exercise 1:** Use NCBI **datasets** to download a genome FASTA file
- *Klebsiella pneumoniae*
- **Exercise 2:** Run **kleborate** on FASTA file for subtyping, serotyping, virulence and AMR prediction



Container Repositories and Writing Dockerfiles

Goals by End of Week Two

- Learn about publicly available container registries & resources
- Understand the Dockerfile and how it is used for building docker images
- Learn best-practices for writing dockerfiles
- Learn how to build a docker image on the command line using pre-defined dockerfiles



Outline

- Container registries & resources
- Getting started with Dockerfiles
 - Best practices for writing Dockerfiles
- Building docker images with **docker build**
- Hands-on exercises with **docker build**

Week 1 Review

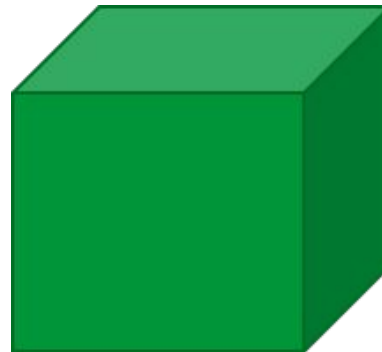
- **Dockerfile** is used to create the docker **image**
- Docker **image** is used to create the docker **container**
- Container **is the runnable instance of an image**

Dockerfile

```
1 FROM ubuntu:xenial
2
3 # metadata
4 LABEL base.image="ubuntu:xenial"
5 LABEL version="1"
6 LABEL software="SPAdes"
7 LABEL software.version="3.13.0"
8 LABEL description="de novo DBG genome assembler"
9 LABEL website="http://cab.spbu.ru/files/release3.13.0/manual.html"
10
11 # Maintainer
12 MAINTAINER Curtis Kapsak <curtis.kapsak@state.co.us>
13
14 RUN apt-get update && apt-get install -y python \
15     wget
16
17 RUN wget http://cab.spbu.ru/files/release3.13.0/SPAdes-3.13.0-Linux.tar.gz && \
18     tar -xzf SPAdes-3.13.0-Linux.tar.gz && \
19     rm -r SPAdes-3.13.0-Linux.tar.gz && \
20     mkdir /data
21
22 ENV PATH="${PATH}:/SPAdes-3.13.0-Linux/bin"
23 WORKDIR /data
```

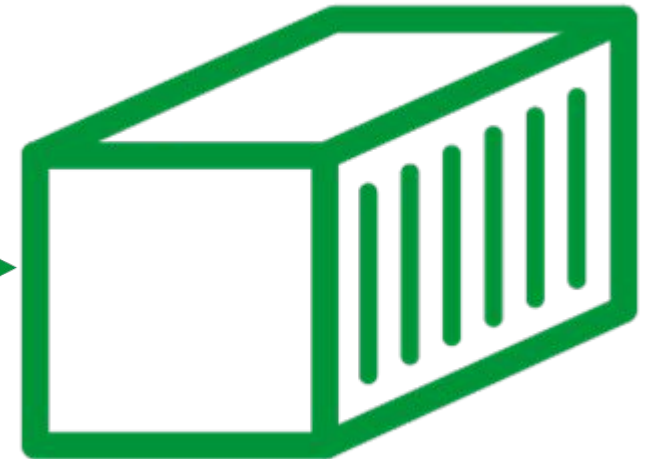
`docker build`

Docker Image



`docker run`

Docker container



Review - Container registries

Docker Images can be built locally **or** pre-built images can be downloaded from public repositories like:

- Docker hub <https://hub.docker.com/>
- Quay.io <https://quay.io/>
- GitHub container registry (GHCR) <https://ghcr.io>
- Cloud provider container registries
 - GCP Artifact Registry
 - Amazon Elastic Container Registry
 - Microsoft Azure Container Registry
- Private registries are an (paid) option



Review - How to share/obtain containers: container registries

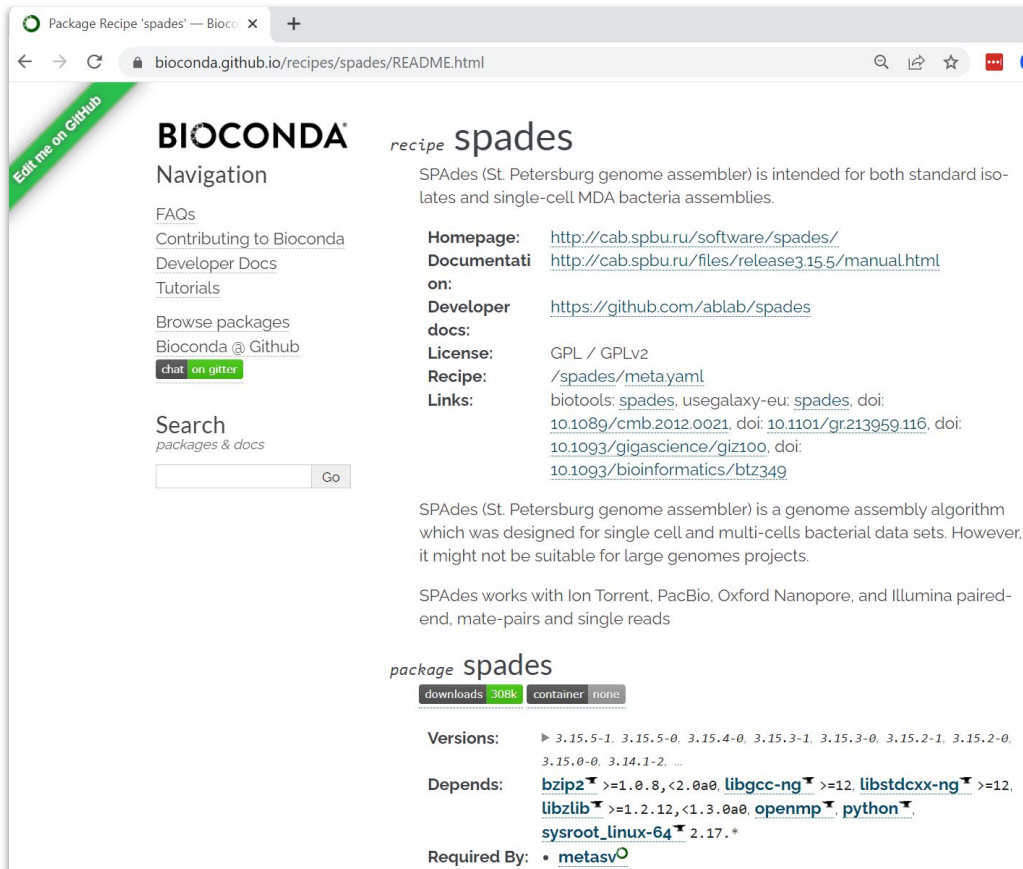
- Docker Hub <https://hub.docker.com/>
 - <https://hub.docker.com/u/staphb>
 - **Millions of public docker images**
- Quay.io
 - <https://quay.io/organization/staphb>
 - (StaPH-B mirrors docker images between dockerhub and quay)



Red Hat
Quay.io

Review - How to share/obtain containers: container registries

- All bioconda packages are available as docker images on quay.io
 - <https://bioconda.github.io/recipes/spades/README.html>



BIOCONDA recipe **spades**

SPAdes (St. Petersburg genome assembler) is intended for both standard isolates and single-cell MDA bacteria assemblies.

Homepage: <http://cab.spbu.ru/software/spades/>
Documentation: <http://cab.spbu.ru/files/release3.15.5/manual.html>
Developer docs: <https://github.com/ablab/spades>
License: GPL / GPLv2
Recipe: /spades/metayaml
Links: biotools: [spades](#), usegalaxy-eu: [spades](#), doi: [10.1089/cmb.2012.0021](#), doi: [10.1101/gr.213959.116](#), doi: [10.1093/gigascience/giz100](#), doi: [10.1093/bioinformatics/btz349](#)

SPAdes (St. Petersburg genome assembler) is a genome assembly algorithm which was designed for single cell and multi-cells bacterial data sets. However, it might not be suitable for large genomes projects.

SPAdes works with Ion Torrent, PacBio, Oxford Nanopore, and Illumina paired-end, mate-pairs and single reads

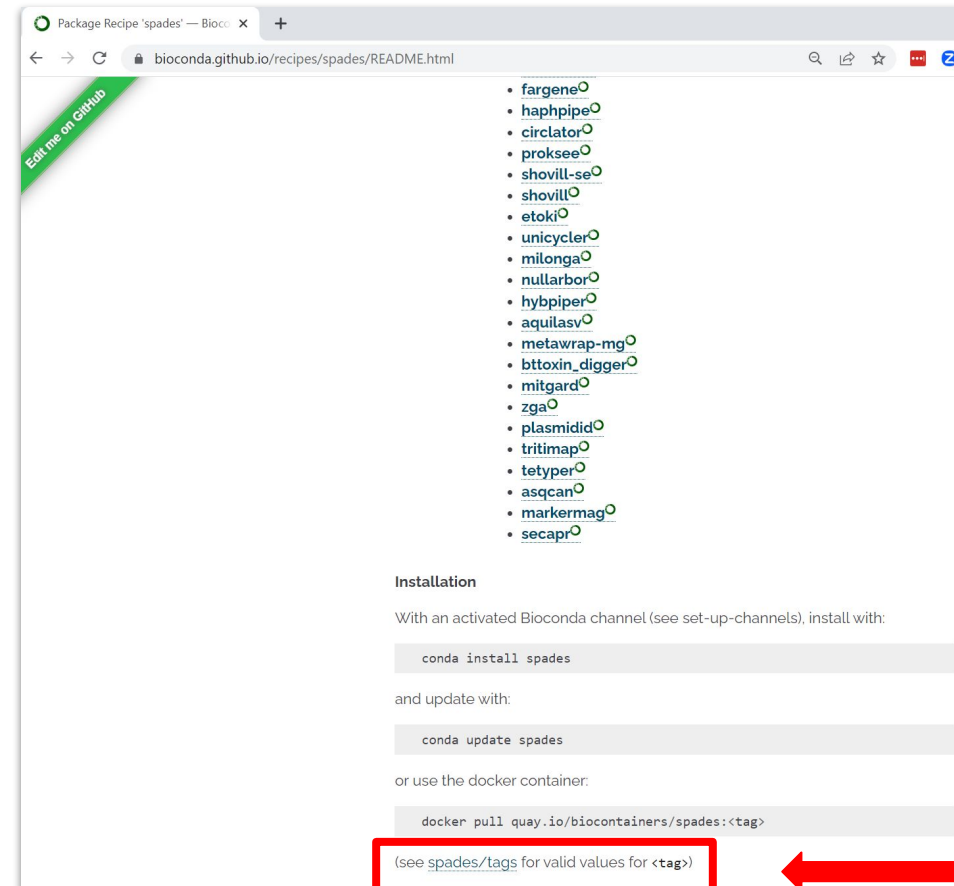
package spades

downloads 308k container none

Versions: 3.15.5-1, 3.15.5-0, 3.15.4-0, 3.15.3-1, 3.15.3-0, 3.15.2-1, 3.15.2-0, 3.15.0-0, 3.14.1-2, ...

Depends: [bzip2](#) >=1.0.8, <2.0a0, [libgcc-ng](#) >=12, [libstdc++-ng](#) >=12, [libzlib](#) >=1.2.12, <1.3.0a0, [openmp](#), [python](#), [sysroot_linux-64](#) 2.17.*

Required By: [metasp](#)



Installation

With an activated Bioconda channel (see set-up-channels), install with:

```
conda install spades
```

and update with:

```
conda update spades
```


or use the docker container:

```
docker pull quay.io/biocontainers/spades:<tag>
```


(see [spades/tags](#) for valid values for <tag>)

link to quay.io


Why write my own dockerfile?

- Not all docker images were created equally
 - Some (like biocontainers) are made by robots! 
- Not all docker images work “out-of-the-box”
 - Limited-to-no testing performed with docker image


Why write my own dockerfile?

- Not all docker images were created equally
 - Some (like biocontainers) are made by robots! 
- Not all docker images work “out-of-the-box”
 - Limited-to-no testing performed with docker image
- “I cannot find a docker image for the software I want to use”

Why write my own dockerfile?

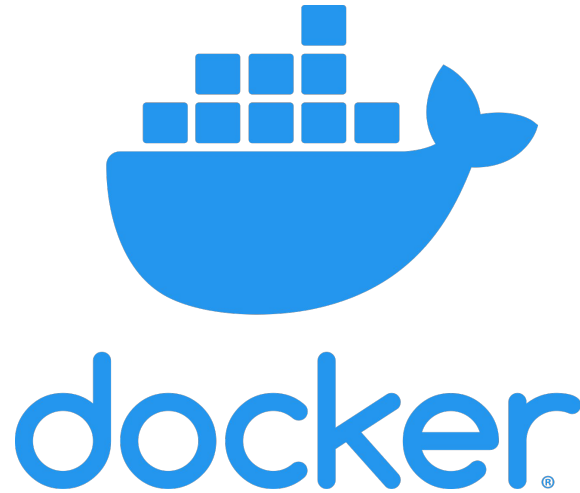
- Not all docker images were created equally
 - Some (like biocontainers) are made by robots! 
- Not all docker images work “out-of-the-box”
 - Limited-to-no testing performed with docker image
- “I cannot find a docker image for the software I want to use”
- “I want to know EXACTLY how X software was installed”
 - “Were versions pinned?”

Why write my own dockerfile?

- Not all docker images were created equally
 - Some (like biocontainers) are made by robots! 
- Not all docker images work “out-of-the-box”
 - Limited-to-no testing performed with docker image
- “I cannot find a docker image for the software I want to use”
- “I want to know EXACTLY how X software was installed”
 - “Were versions pinned?”
- “I want to include multiple tools in a single docker image”
 - minimap2 + samtools commands piped together:

```
minimap2 -x map-ont -a ref.fasta reads.fastq.gz | samtools sort -o out.bam
```

Getting started with Dockerfiles



Review - The Dockerfile

- In order to build a docker image, you need at a minimum one file: the **Dockerfile**
- **Dockerfile** = set of instructions used to build a docker image
- Similar to an installation script or a **.yaml** file used for making/sharing conda environments

Format

Here is the format of the `Dockerfile` :

```
# Comment  
INSTRUCTION arguments
```

<https://docs.docker.com/engine/reference/builder/>

```
FROM ubuntu:jammy
```

```
# install stuff!
```

```
RUN install \  
    softwareA \  
    softwareB \  
    softwareC
```

The Dockerfile

- Dockerfile instructions (**FROM**, **RUN**, **COPY**, **ENV**, etc.) will add a “layer” to the docker image
- Images are multi-layered and different images may share layers like the base image
 - **FROM** ubuntu:focal

Spades Dockerfile

- <https://github.com/StaPH-B/docker-builds/blob/master/spades/3.15.5/Dockerfile>

Dockerfile Cheat Sheet

- https://kapeli.com/cheat_sheets/Dockerfile.docset/Contents/Resources/Documents/index

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```


Dockerfile - FROM

- Initializes a new build stage
- Required - A valid Dockerfile must start with a **FROM** instruction
 - The only instruction that can precede **FROM** is an **ARG** variable (more on this later)
- **FROM** defines the base image
 - Recommendation - choose a base image and stick with it
- Official docs:
<https://docs.docker.com/engine/reference/builder/#from>

FROM




```
FROM [--platform=<platform>] <image> [AS <name>]
```



Or

```
FROM [--platform=<platform>] <image>[:<tag>] [AS <name>]
```

Or

```
FROM [--platform=<platform>] <image>[@<digest>] [AS <name>]
```

 master  docker-builds / spades / 3.15.5 / Dockerfile 

 kapsakcj added ca-certificates to spades 3.15.5 dockerfile 

Code

Blame42 lines (34 loc) · 1.39 KB

1 FROM ubuntu:focal as app

Dockerfile - FROM

My favorite base images

Ubuntu - hub.docker.com/_/ubuntu

- Familiar linux OS - has many basic linux commands installed (**ls**, **cd**, **cp**, **mv**, **ps**, etc.)
- Relatively easy to install dependencies via **apt-get**
 - packages.ubuntu.com for looking up what is available via **apt-get**
- Can use **pip** for python packages & **cpan/cpanm** for perl packages

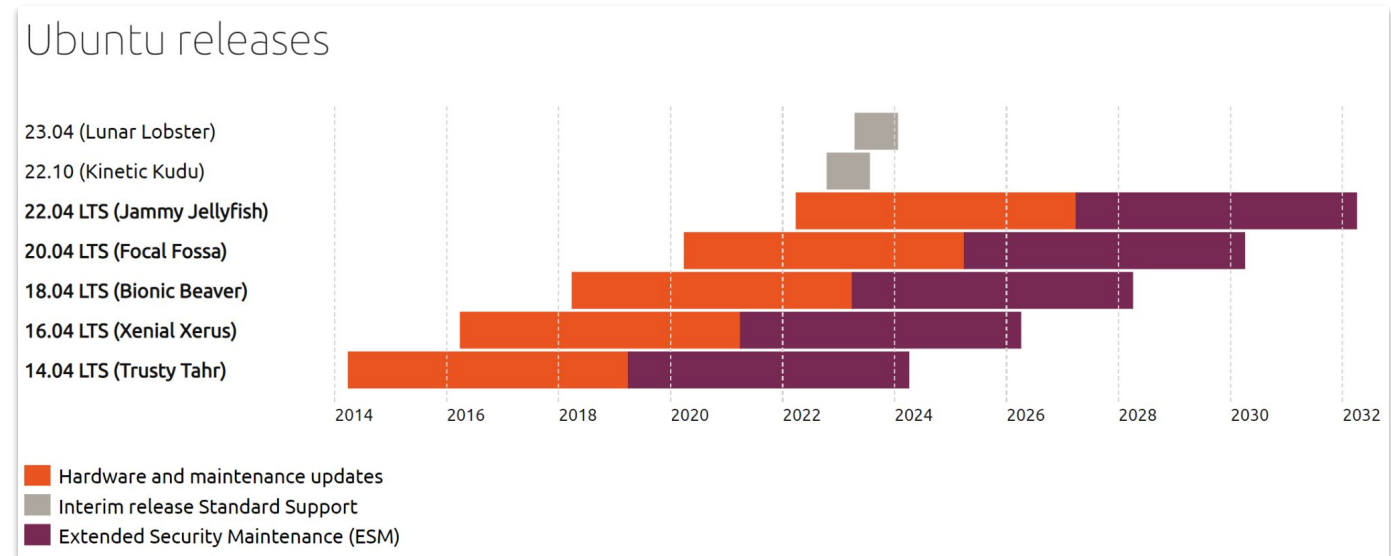
```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

Dockerfile - FROM

My favorite base images

Ubuntu - hub.docker.com/_/ubuntu

- Recommendation
 - Ubuntu 20.04 LTS (focal)
 - **FROM** `ubuntu:focal`
 - Ubuntu 22.04 LTS (jammy)
 - **FROM** `ubuntu:jammy`
 - Or the next LTS (Long Term Support) release
- Older Ubuntu releases are nearing the end of support, use something new that will be supported long term!



<https://ubuntu.com/about/release-cycle#ubuntu>

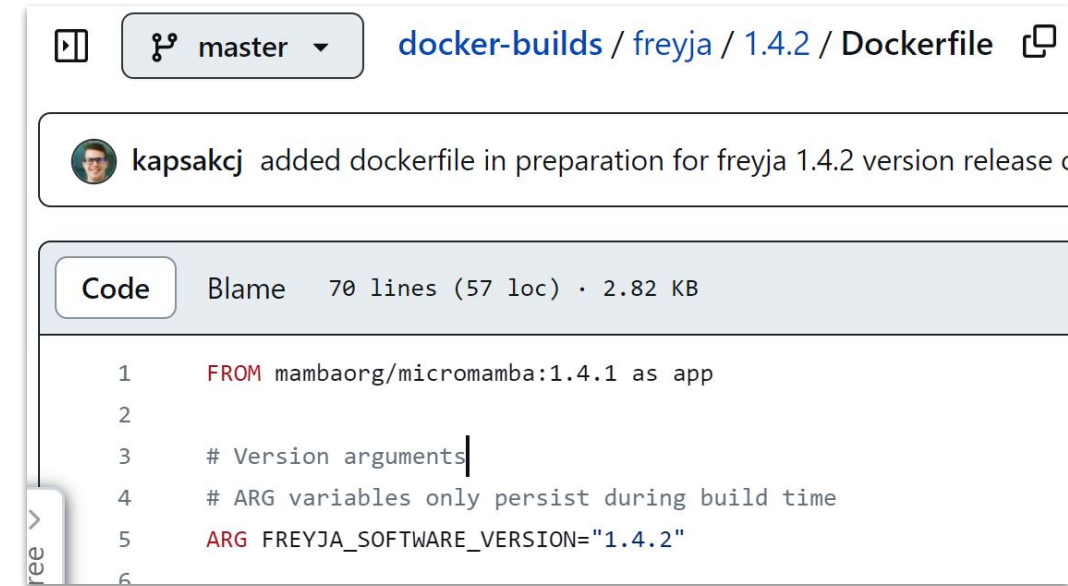
Dockerfile - FROM

My favorite base images:

micromamba -

hub.docker.com/r/mambaorg/micromamba

- Familiar linux OS (Debian)
- **micromamba** is preinstalled
 - **micromamba** is even more lightweight than conda or miniconda
- I use for complicated installations where I rely upon the conda package in bioconda
← not best practice!
- Also use for scenarios where I want to use a conda recipe file (.yml) for installation



The screenshot shows a GitHub repository for 'docker-builds / freyja / 1.4.2 / Dockerfile'. A commit by 'kapsakcj' is shown, adding a Dockerfile. The 'Code' tab is selected, showing a Dockerfile with 70 lines (57 loc) and 2.82 KB. The visible code includes:

```
1 FROM mambaorg/micromamba:1.4.1 as app
2
3 # Version arguments
4 # ARG variables only persist during build time
5 ARG FREYJA_SOFTWARE_VERSION="1.4.2"
6
```

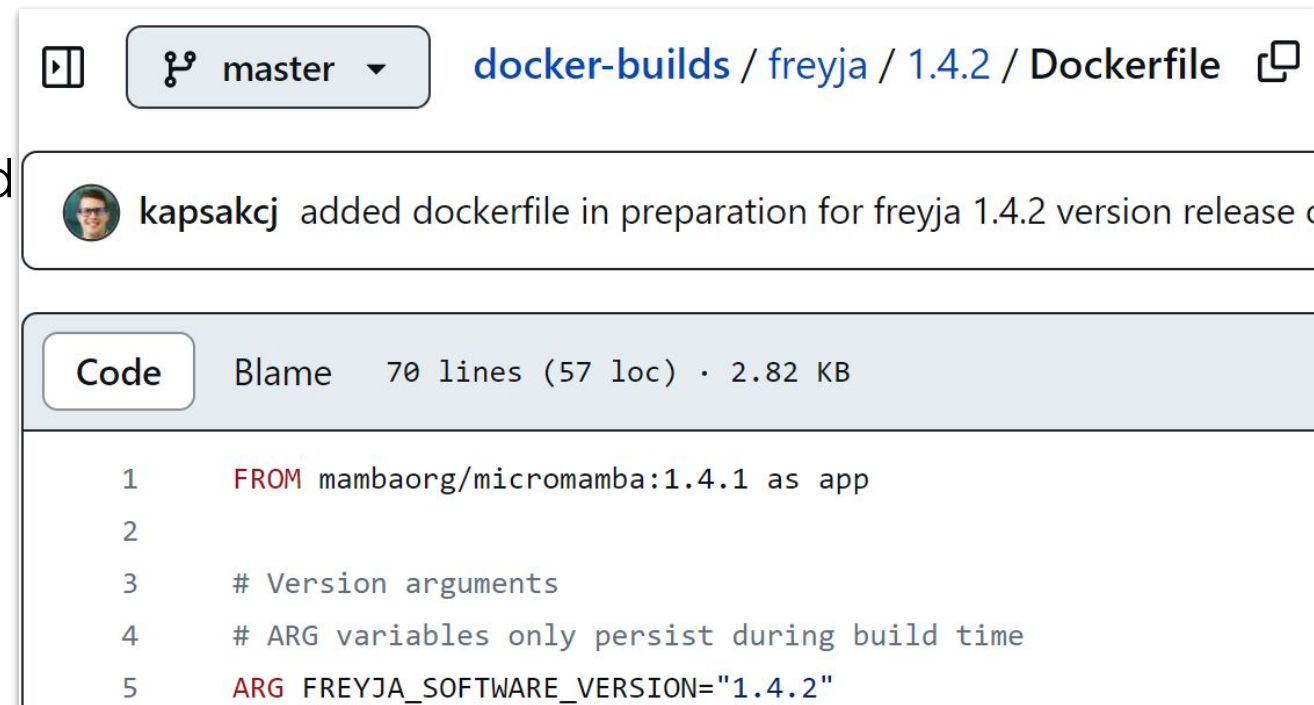
later in Freyja dockerfile 

```
33 # Create Freyja conda environment called freyja-env from bioconda recipe
34 # clean up conda garbage
35 RUN micromamba create -n freyja-env -c conda-forge -c bioconda -c defaults freyja=${FREYJA_SOFTWARE_VERSION} && \
36     micromamba clean -a -y
```


Dockerfile - ARG

- Sets environmental variables that are ONLY available during docker image build time
- Once image is built, all **ARG** variables are unset/removed
- Useful for specifying versions of tools to install; can make it easy to upgrade versions by only changing one line of code

beginning of Freyja dockerfile 



```
1  FROM mambaorg/micromamba:1.4.1 as app
2
3  # Version arguments
4  # ARG variables only persist during build time
5  ARG FREYJA_SOFTWARE_VERSION="1.4.2"
```

later in Freyja dockerfile 

```
33  # Create Freyja conda environment called freyja-env from bioconda recipe
34  # clean up conda garbage
35  RUN micromamba create -n freyja-env -c conda-forge -c bioconda -c defaults freyja=${FREYJA_SOFTWARE_VERSION} && \
36  micromamba clean -a -y
```

Dockerfile - **ENV**

- Sets permanent environmental variables that are available during image build and afterwards for all users
- Useful for setting the **\$PATH** variable
- [Format:](#)

ENV 

```
ENV <key>=<value> ...
```

SPAdes dockerfile example:

```
32     # set PATH and locale settings for singularity
33     ENV LC_ALL=C.UTF-8 \
34         PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
```

- Can set multiple variables in one **ENV** layer using line breaks with a backslash \
- Useful for tools that do not automatically get added to your **\$PATH** variable

Dockerfile - RUN

- Executes a command in a new layer
- Each **RUN** layer builds upon the previous **FROM** and **RUN** layers
- Changes made in **RUN** commands are saved in the final docker image
- Assume **/bin/sh** shell for running commands
- Docker official docs: <https://docs.docker.com/engine/reference/builder/#run>

RUN

RUN has 2 forms:

- `RUN <command>` (*shell* form, the command is run in a shell, which by default is `/bin/sh -c` on Linux or `cmd /S /C` on Windows)
- `RUN ["executable", "param1", "param2"]` (*exec* form)

The `RUN` instruction will execute any commands in a new layer on top of the current image and commit the results. The resulting committed image will be used for the next step in the `Dockerfile`.

Dockerfile - RUN

- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
 - `apt-get update && apt-get install` in a single **RUN** statement, done early in dockerfile



```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```


Dockerfile - RUN


- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
 - **apt-get update && apt-get install** in a single **RUN** statement, done early in dockerfile
 - Use syntax appropriate for your build stage (shell in ubuntu). Imagine you are using the command line



```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

Dockerfile - RUN

- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
 - **apt-get update && apt-get install** in a single **RUN** statement, done early in dockerfile
 - Use syntax appropriate for your build stage (shell in ubuntu). Imagine you are using the command line
 - \ (backslashes) are for line breaks (readability)



```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

Dockerfile - RUN


- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
 - **apt-get update && apt-get install** in a single **RUN** statement, done early in dockerfile
 - Use syntax appropriate for your build stage (shell in ubuntu). Imagine you are using the command line
 - \ (backslashes) are for line breaks (readability)
 - **&&** bash operator is used to create long one-liners so that multiple commands are run sequentially and are dependent on each command running successfully



```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

Dockerfile - RUN

- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
 - **apt-get update && apt-get install** in a single **RUN** statement, done early in dockerfile
 - Use syntax appropriate for your build stage (shell in ubuntu). Imagine you are using the command line
 - \ (backslashes) are for line breaks (readability)
 - **&&** bash operator is used to create long one-liners so that multiple commands are run sequentially and are dependent on each command running successfully
 - Assume **/bin/sh** shell for running commands, but there is a way to set **`/bin/bash`** as normal shell. Not necessary unless using bash-specific cmdline tricks



```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

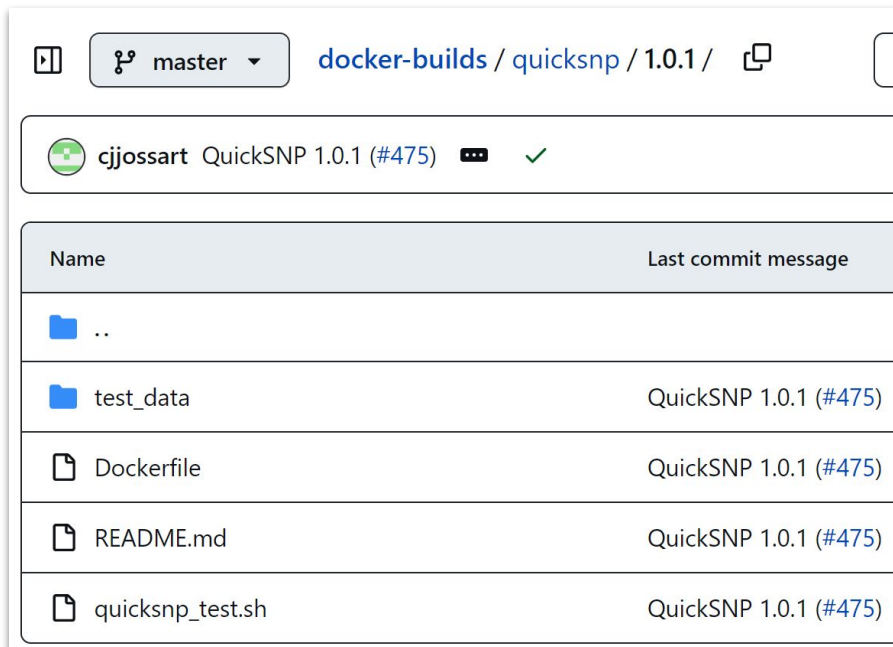

Dockerfile - WORKDIR

- Sets the working directory for any **RUN**, **CMD**, **ENTRYPOINT**, **COPY** and **ADD** instructions that follow it in the Dockerfile.
- **WORKDIR** will also persist after the image is build
- Can use multiple **WORKDIR**'s in one dockerfile, but only last **WORKDIR** will be saved in final image
- All StaPH-B/docker-builds images have the final **WORKDIR** set to **/data**
 - used for passing data in and out of containers
 - does not overlap with existing directories in container filesystem

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

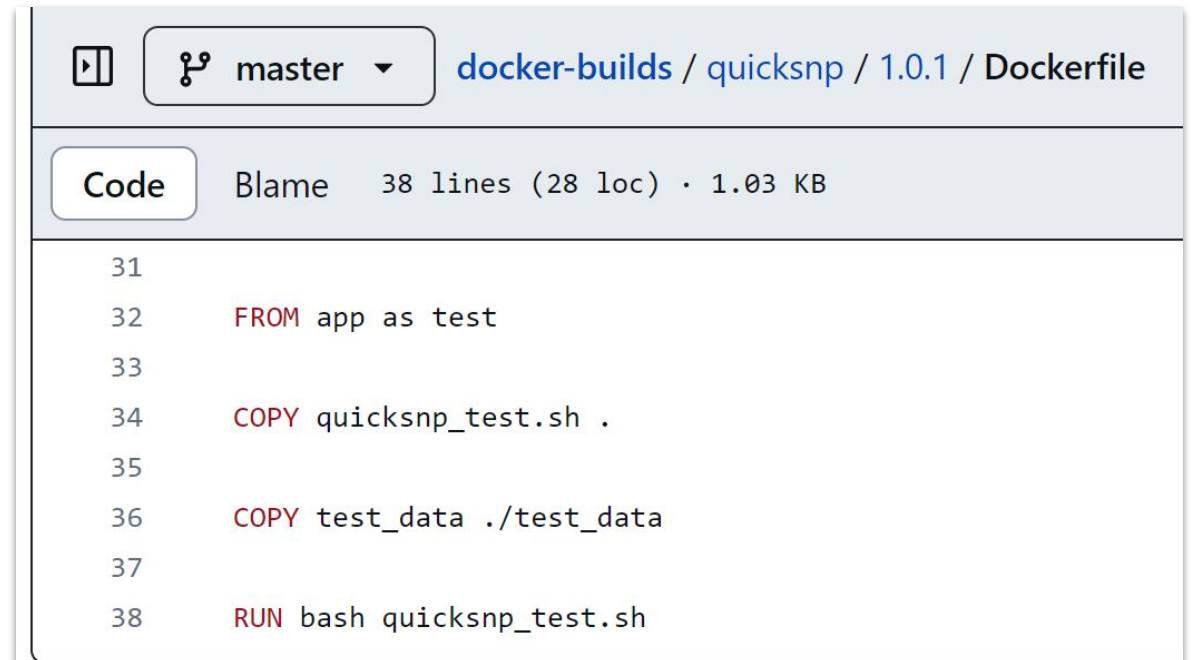
Dockerfile - COPY

- Copies files into the image
- Files must be either:
 - located in same directory as Dockerfile (AKA the “build context”)
 - From another stage of the build process
- Useful for copying in test data, (small) databases, additional code:



The screenshot shows a GitHub repository interface. At the top, the breadcrumb is 'docker-builds / quicksnip / 1.0.1'. Below it, the repository name 'cjjossart QuickSNIP 1.0.1 (#475)' is displayed with a green checkmark. A table lists the files in the directory:

Name	Last commit message
..	
test_data	QuickSNIP 1.0.1 (#475)
Dockerfile	QuickSNIP 1.0.1 (#475)
README.md	QuickSNIP 1.0.1 (#475)
quicksnip_test.sh	QuickSNIP 1.0.1 (#475)



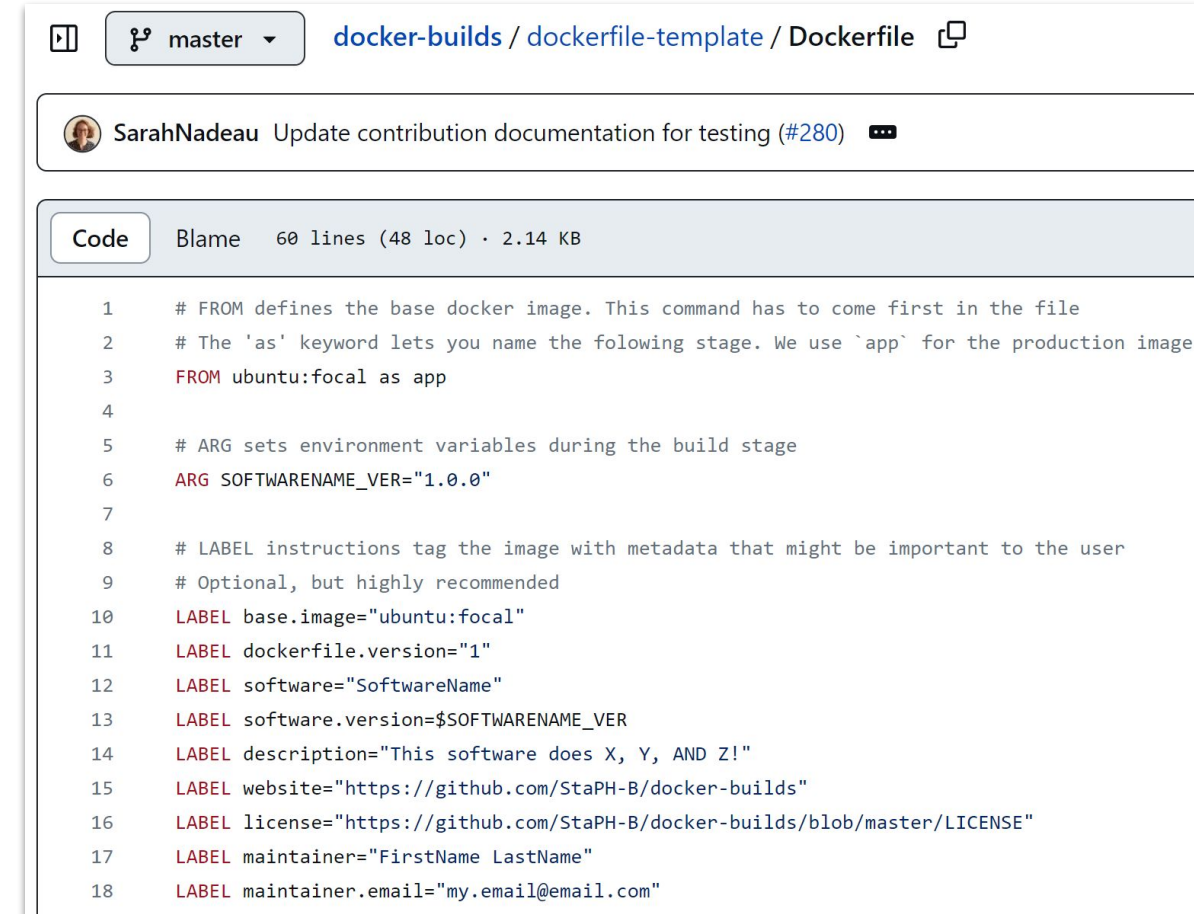
The screenshot shows the Dockerfile content for the repository. The breadcrumb is 'docker-builds / quicksnip / 1.0.1 / Dockerfile'. The file is 38 lines (28 loc) and 1.03 KB. The content is as follows:

```
31
32 FROM app as test
33
34 COPY quicksnip_test.sh .
35
36 COPY test_data ./test_data
37
38 RUN bash quicksnip_test.sh
```

Dockerfile - LABEL

- optional but highly recommended!
- allows addition of metadata to your docker image
- generally located near top of dockerfile, after **FROM**
- StaPH-B has some required labels for dockerfiles contributed to their repo
 - See the template Dockerfile for examples:

<https://github.com/StaPH-B/docker-builds/blob/master/dockerfile-template/Dockerfile>

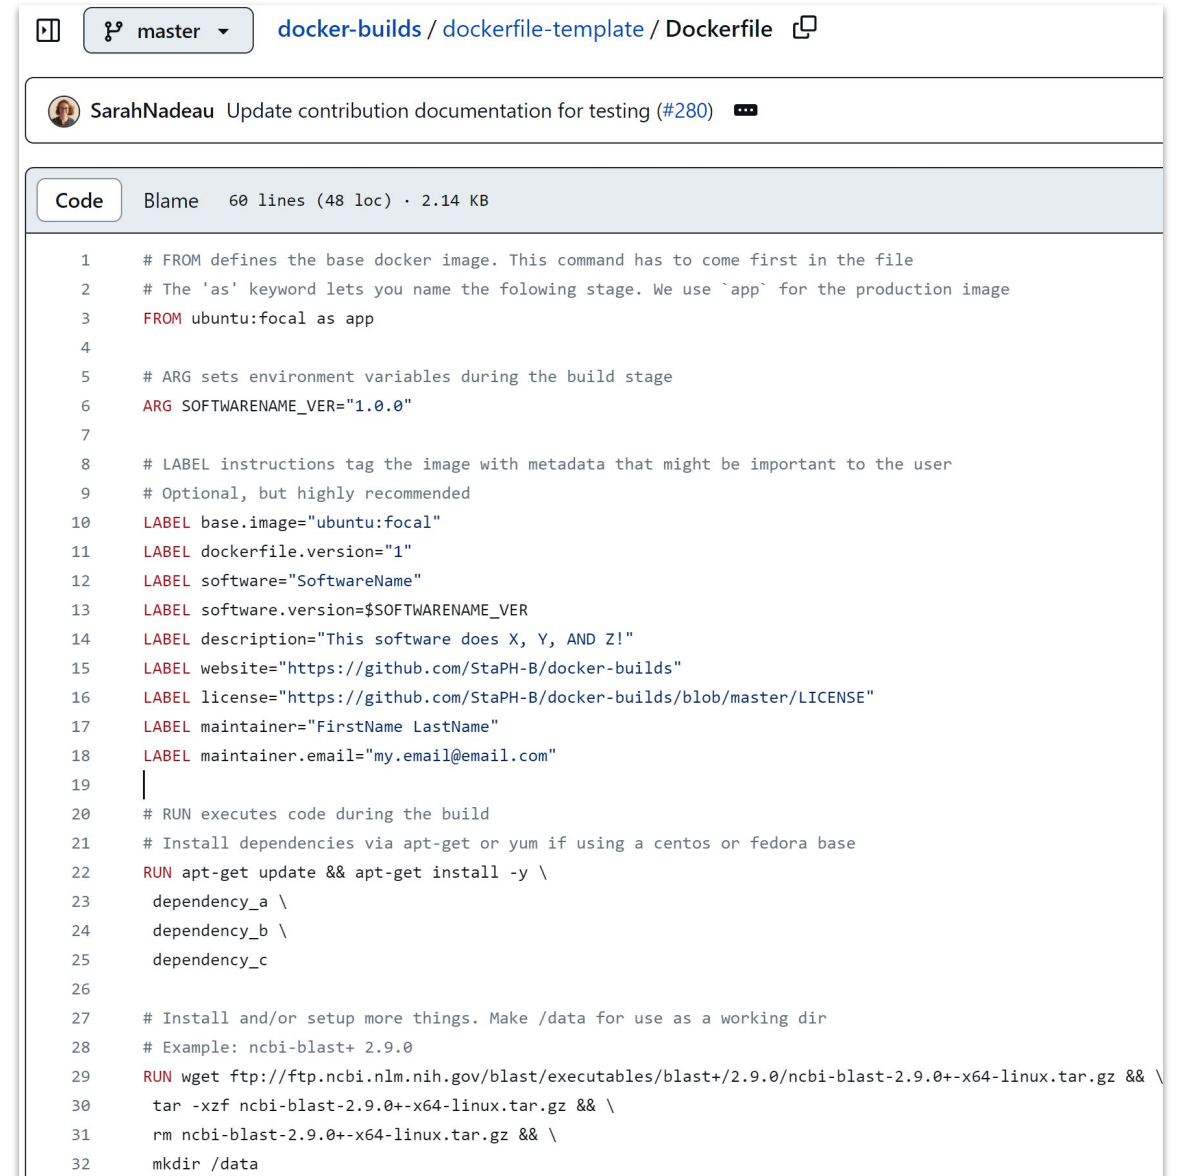


The screenshot shows a GitHub repository interface for the 'docker-builds' repository, specifically the 'dockerfile-template' directory. The file 'Dockerfile' is selected, showing its content. The interface includes a header with the repository name and a dropdown menu for branches (currently set to 'master'). Below the header, there is a commit message by SarahNadeau: 'Update contribution documentation for testing (#280)'. The main area displays the Dockerfile code, which includes comments and instructions for setting the base image, environment variables, and labels.

```
1  # FROM defines the base docker image. This command has to come first in the file
2  # The 'as' keyword lets you name the following stage. We use `app` for the production image
3  FROM ubuntu:focal as app
4
5  # ARG sets environment variables during the build stage
6  ARG SOFTWARENAME_VER="1.0.0"
7
8  # LABEL instructions tag the image with metadata that might be important to the user
9  # Optional, but highly recommended
10 LABEL base.image="ubuntu:focal"
11 LABEL dockerfile.version="1"
12 LABEL software="SoftwareName"
13 LABEL software.version=$SOFTWARENAME_VER
14 LABEL description="This software does X, Y, AND Z!"
15 LABEL website="https://github.com/StaPH-B/docker-builds"
16 LABEL license="https://github.com/StaPH-B/docker-builds/blob/master/LICENSE"
17 LABEL maintainer="FirstName LastName"
18 LABEL maintainer.email="my.email@email.com"
```

Dockerfile - comments

- also optional but highly recommended!
- comment lines begin with **#**



The screenshot shows a GitHub repository for 'docker-builds / dockerfile-template / Dockerfile'. The file is 60 lines long (48 loc) and 2.14 KB. It contains a Dockerfile with the following content:

```
1  # FROM defines the base docker image. This command has to come first in the file
2  # The 'as' keyword lets you name the following stage. We use `app` for the production image
3  FROM ubuntu:focal as app
4
5  # ARG sets environment variables during the build stage
6  ARG SOFTWARENAME_VER="1.0.0"
7
8  # LABEL instructions tag the image with metadata that might be important to the user
9  # Optional, but highly recommended
10 LABEL base.image="ubuntu:focal"
11 LABEL dockerfile.version="1"
12 LABEL software="SoftwareName"
13 LABEL software.version=$SOFTWARENAME_VER
14 LABEL description="This software does X, Y, AND Z!"
15 LABEL website="https://github.com/StaPH-B/docker-builds"
16 LABEL license="https://github.com/StaPH-B/docker-builds/blob/master/LICENSE"
17 LABEL maintainer="FirstName LastName"
18 LABEL maintainer.email="my.email@email.com"
19 |
20 # RUN executes code during the build
21 # Install dependencies via apt-get or yum if using a centos or fedora base
22 RUN apt-get update && apt-get install -y \
23     dependency_a \
24     dependency_b \
25     dependency_c
26
27 # Install and/or setup more things. Make /data for use as a working dir
28 # Example: ncbi-blast+ 2.9.0
29 RUN wget ftp://ftp.ncbi.nlm.nih.gov/blast/executables/blast+/2.9.0/ncbi-blast-2.9.0+-x64-linux.tar.gz && \
30     tar -xzf ncbi-blast-2.9.0+-x64-linux.tar.gz && \
31     rm ncbi-blast-2.9.0+-x64-linux.tar.gz && \
32     mkdir /data
```


Break

5 min break!
resume at 3:38pm

Quiz time:

- What is the base image used in iVar 1.4.1 dockerfile?
 - <https://github.com/StaPH-B/docker-builds/blob/master/ivar/1.4.2/Dockerfile>
- What **ARG** variables are used and what are they used for?
- How is `ivar` installed? What steps are taken?

Dockerfile → Docker image

docker build

Docker Build

- Builds an image from a dockerfile
- At a minimum, requires a Dockerfile. Some dockerfiles require other files for building (scripts, databases, etc.)
- Official docs: <https://docs.docker.com/engine/reference/commandline/build/>
- General command structure:

`docker build --tag <name>:<tag> <directory-with-dockerfile>`

- example using SPAdes dockerfile:

`docker build --tag spades:3.15.5 spades/3.15.5/`

Multi-stage docker builds

- Docker images that use multiple **FROM** instructions
 - can use a new base image **OR**
 - can use previous stage
- Useful for “optimizing” Dockerfiles.
Optimizing =
 - reducing number of layers
 - removing unnecessary software & files, etc. to reduce final size of docker image
- Not required in general, but are required for StaPH-B/docker-builds project

General format in Dockerfile (using previous stage as new **FROM** layer):

```
...
```

```
FROM baseimage as name
```

```
RUN stuff
```

```
FROM name as name2
```

```
RUN stuff
```

```
FROM name2 as final
```

```
RUN stuff
```

```
...
```

Multi-stage docker builds

General format in Dockerfile (using previous stage as new FROM layer):

...

```
FROM ubuntu:focal as stage1
```

```
RUN stuff
```

```
FROM stage1 as stage2
```

```
RUN stuff
```

```
FROM stage2 as finalstage
```

```
RUN stuff
```

...

Multi-stage docker builds

General format in Dockerfile (using new base images as new FROM layer):

...

```
FROM ubuntu:focal as stage1
```

```
RUN stuff
```

```
FROM python:slim as stage2
```

```
RUN stuff
```

```
FROM ubuntu:jammy as finalstage
```

```
RUN stuff
```

...

Multi-stage docker builds

- StaPH-B exclusive
- Require use of 2 stages, **app** and **test**
- **app** stage
 - stage for installing software and dependencies
 - only stage that remains in final docker image
- **test** stage
 - based on **app** stage
 - runs tests to ensure functionality and correct version of installed software
 - stage is NOT included in final docker image

General format in StaPH-B Dockerfiles:

...

```
FROM ubuntu:focal as app
RUN install software
```

app stage

```
FROM app as test
```

```
RUN my-tests.sh
```

```
RUN software --version
```

...

test stage

Multi-stage docker builds

- Let's switch to GitPod and see this in practice
- Navigate to the training exercise document:
 - <https://github.com/theiagen/docker-builds/tree/master/training/NE-BRR-docker-for-PH-bioinformatics-May2023>
- Sign into Gitpod:
 - <https://gitpod.io/workspaces>

Further reading & resources

- StaPH-B Github repo and docker hub account
 - <https://github.com/StaPH-B/docker-builds>
 - <https://hub.docker.com/u/staphb>
- Docker Documentation - a wealth of info here. Note that we use Docker Community Edition, as you have to pay for the Enterprise Edition
 - <https://docs.docker.com/>
- An awesome tutorial/workshop on docker for bioinformatics
 - <https://github.com/PawseySC/bio-workshop-18>
- Template for your Dockerfile
 - <https://github.com/StaPH-B/docker-builds/blob/master/dockerfile-template/Dockerfile>
- Some best practices
 - https://staphb.org/docker-builds/make_containers/
- Search for docker images and (sometimes) Dockerfiles here:
 - <http://hub.docker.com/>
 - <https://quay.io/>
- “What is Docker?” (~11 min)
 - https://www.youtube.com/watch?time_continue=1&v=aLipr7tTuA4

Acknowledgements

- MA DPH
- Members of StaPH-B & the docker-builds contributors & maintainers
 - Erin Young, UT PHL
 - Kelsey Florek, WI PHL
 - Kevin Libuit, Theiagen Genomics
 - Frank Ambrosio, Theiagen Genomics
 - many more awesome people!
 - StaPH-B docker-builds contributors:
<https://github.com/StaPH-B/docker-builds#authorsmaintainers>
- APHL
- CDC

